

Highlights

- Human task scheduling is cast as a POMDP with latent attention and fatigue.
- Rest breaks are explicit actions chosen by a memory-augmented deep Q-network.
- Agent beats tuned break heuristics: 72.8% vs 71.4% of tasks on time.
- The learned policy rests during the post-lunch dip and does hard tasks early.
- Ablations show the gain depends on modelling dynamic human cognition.