

Highlights

- Human task scheduling is cast as a POMDP with latent attention and fatigue.
- Rest breaks are explicit actions of a memory-augmented deep Q-network.
- Tuned rest rules are strong; learned agents match or beat the best of nine.
- Only learned policies protect hard tasks while maximising on-time completion.
- The value of learning grows sharply with the accuracy of cognitive sensing.