

Supplementary Material

Attention-aware task scheduling under partial observability of human cognitive state: a memory-augmented deep reinforcement learning approach

P. N. Bhargav Teja, Lanka Sree Chathurya, Rajay Vedaraj I.S.*
Vellore Institute of Technology, Vellore 632014, Tamil Nadu, India

Contents

S1 Complete simulator specification	1
S1.1 Episode timeline	1
S1.2 Productivity function	2
S1.3 Environment variants used in the ablations	2
S2 Reproducibility	2
S3 Baseline tuning	3
S4 Checkpoint selection	3
S5 Component and sensitivity analysis	3
S6 Audit trail of the baseline set and action space	3
S7 Full robustness results	4
S8 Reward-weight sweep	4

S1 Complete simulator specification

S1.1 Episode timeline

Each episode (workday) is generated from a single integer seed that fixes, through independent draws from one pseudo-random generator: the arrival slots, difficulties and slack of all tasks; the initial state (H_0, F_0); the process-noise sequence $\xi_{0:T-1}$; and the observation-noise sequence $\epsilon_{0:T}$. Because the noise sequences are pre-drawn, two policies evaluated on the same seed experience the same exogenous randomness (common random numbers), and their outcomes differ only through their decisions. Within slot t :

1. Tasks with $a_i \leq t$ become available.
2. The scheduler receives o_t (main text, Eq. 6), computed from A_t, F_t and ϵ_t .
3. The action is mapped to a task (EDF within the chosen class, with a fallback to the other class) or to rest.

*Corresponding author: rajay@vit.ac.in

4. If working on task i : the remaining work is decreased by ρ_t (Eq. 5); if it reaches zero, the task is completed at $t + 1$ and the task reward is credited; H and F are updated with the work branch of Eqs. (2)–(3). If resting: the rest branch is applied.
5. Process noise is added to H , and both states are clipped to $[0, 1]$.
6. The cognitive reward $\mu A_{t+1} - \lambda F_{t+1}$ is added, and every unfinished task whose deadline equals $t + 1$ incurs the miss penalty (once).

S1.2 Productivity function

Figure S1 shows the work rate $\rho = \rho_0 A^{0.5+d}(1 - \kappa F)$ for representative difficulties and fatigue levels. At $A = 0.9$ and $F = 0.1$, an easy task ($d = 0.2$) proceeds at 0.88 units/slot and a hard task ($d = 1$) at 0.81. At $A = 0.4$ and $F = 0.8$, the rates fall to 0.32 and 0.15, respectively. Hard work is thus disproportionately penalised in low-attention states, which creates the incentive to match task difficulty to cognitive state.




Figure S1: Work rate as a function of attention for three difficulty levels and two fatigue levels.

S1.3 Environment variants used in the ablations

- **No observation noise:** $\sigma_o = 0$; the cognitive state is observed exactly (an MDP with respect to the cognitive variables, although future arrivals remain hidden).
- **No circadian process:** $\gamma = 0$, so that $A_t = H_t$.
- **Static capacity (no cognition):** $A_t \equiv 0.60$ and $F_t \equiv 0.30$ at all times. The resulting work rate of $0.60^{0.5+d} \times 0.85$ (0.48 units/slot for $d = 0.6$) delivers a daily throughput comparable to the break-taking heuristics in the full model, so the variant removes *dynamics* without drastically changing *capacity*. Breaks have no benefit in this variant.
- **Task-only reward:** $\lambda = \mu = 0$; the dynamics are unchanged.
- **Worker profiles:** resilient ($\alpha, \delta \times 0.7$; $\beta, \eta \times 1.3$) and fatigue-prone ($\alpha, \delta \times 1.3$; $\beta, \eta \times 0.75$).
- **Workload:** $N = 8$ (4 initial) and $N = 16$ (8 initial) tasks per day.

S2 Reproducibility

All experiments can be reproduced with the released code:

```
pip install -r requirements.txt
python -m pytest # simulator unit tests
python -m experiments.run_eval heuristics # tune + test rule-based baselines
```

```
python -m experiments.run_rl --workers 4 # train all 107 agents
python -m experiments.run_eval robustness # zero-shot robustness study
python -m experiments.run_eval traces # behavioural traces
python -m analysis.analyze # all tables, figures, statistics
```

Seed ranges: training days = $s \times 1\,000\,003 + e$ for training seed s and episode e ; validation days $10^7, \dots, 10^7 + 99$; test days $2 \times 10^7, \dots, 2 \times 10^7 + 499$; tuning days $3 \times 10^7, \dots, 3 \times 10^7 + 299$. One training run takes about 720 s (AA-DQN) on a single CPU thread.

S3 Baseline tuning

Table S1 lists the parameters selected by grid search for each environment variant (maximum mean return on 300 tuning days). “off” indicates that the fatigue threshold was never triggered ($\tau_F = 1.01$). In the static-capacity variant, rest has no benefit, and all threshold rules learn to (almost) never rest.

Table S1: Tuned heuristic parameters per environment variant.

Variant	EDF-Periodic (p)	EDF-Threshold (τ_F, τ_A)	Cog-Heuristic (τ_F, τ_A, τ_H)
Full model	4	(0.6, 0.7)	(0.4, 0.6, 0.9)
No observation noise	4	(0.3, 0.7)	(0.3, 0.6, 0.9)
No circadian process	4	(0.4, 0.7)	(0.4, 0.6, 0.9)
Static capacity	10	(0.7, 0)	(0.7, 0, 0.8)
Task-only reward	4	(0.7, 0.7)	(0.5, 0.6, 0.9)
$\lambda = 0.05$	4	(0.6, 0.7)	(0.4, 0.6, 0.9)
$\lambda = 0.2$	3	(0.4, 0.7)	(0.3, 0.6, 0.9)
$\lambda = 0.4$	2	(0.3, 0.7)	(0.3, 0.6, 0.9)

S4 Checkpoint selection

Table S2 compares, for each learned policy, the validation-selected checkpoint (primary results) with the final iterate after 5000 episodes. Selection improves both agents, which is consistent with the policy oscillation visible in Fig. S2. The ranking AA-DQN > DQN holds for both checkpoint choices.

Table S2: Test performance (500 days) of validation-selected versus final-iterate checkpoints (mean $\pm$ s.d. over ten seeds).

Method	Checkpoint	Return	On-time rate	Mean selected episode
DQN	validation-selected	10.29 ± 0.26	0.677 ± 0.008	3490
DQN	final iterate	9.84 ± 0.34	0.663 ± 0.011	3000
AA-DQN	validation-selected	12.17 ± 0.15	0.728 ± 0.005	3350
AA-DQN	final iterate	11.55 ± 0.35	0.710 ± 0.013	3000

S5 Component and sensitivity analysis

Table S3 reproduces the component study of the main text with all configurations, including the history length k and the recurrent (GRU) encoder.

S6 Audit trail of the baseline set and action space

For transparency, we document a design change made during the study. In the first version of the experiments, the agent’s action set contained only the intensity actions {EASY, HARD, BREAK} of the original formulation,




Figure S2: Training return (left; 50-episode moving average, ϵ -greedy) and Huber TD loss (right) for DQN and AA-DQN (mean $\pm$ s.d. over ten seeds).

Table S3: Component ablation and sensitivity (mean $\pm$ s.d. over seeds; 500 test days).

Configuration	Seeds	Return	On-time rate	Train time (s)
AA-DQN (history $k=4$ + Double + Dueling)	10	12.19 ± 0.27	0.726 ± 0.009	474
Vanilla DQN ($k=1$, no Double/Dueling)	10	10.66 ± 0.33	0.686 ± 0.013	410

with EDF dispatch within the chosen class. The rule-based baselines were EDF-based (EDF-Periodic, EDF-Threshold, Cog-Heuristic). A pre-submission review argued that SPT-based rest heuristics should also be included, because SPT was the best classical rule in the static-capacity ablation. When added, SPT-Threshold obtained a higher test return than the intensity-only agent. The reason was structural: that agent could not express SPT sequencing. We therefore adopted the hybrid action set, in which every baseline is representable, added the observation of each class’s shortest remaining work, set the training budget to 5000 episodes (the upper end of the range in our original protocol), and re-ran all experiments. To avoid over-fitting the design to the test set, this decision was made once, and no further design changes were made after the final runs were launched. The intensity-only agent trained under the final protocol is reported as a component ablation (Table S3).

S7 Full robustness results

Table S4 reports the return and on-time rate of every method under each zero-shot test condition. Learned policies are averaged over training seeds, and heuristics use the parameters tuned in the nominal environment.

Table S4: Zero-shot robustness: return / on-time rate on 500 test days per condition. Bold marks the best return.

Test condition	AA-DQN	DQN	EDF-Periodic	EDF-Threshold	Cog-Heuristic	EDF
$\sigma_o = 0$	12.11 / 0.725	10.09 / 0.669	11.41 / 0.660	11.00 / 0.655	9.82 / 0.624	4.24 / 0.468
Nominal ($\sigma_o = 0.1$)	12.19 / 0.726	10.66 / 0.686	11.41 / 0.660	11.35 / 0.654	10.59 / 0.638	4.24 / 0.468
$\sigma_o = 0.2$	12.32 / 0.728	11.22 / 0.701	11.41 / 0.660	10.77 / 0.615	10.01 / 0.587	4.24 / 0.468
$\sigma_o = 0.3$	12.40 / 0.725	11.66 / 0.711	11.41 / 0.660	9.63 / 0.557	8.04 / 0.482	4.24 / 0.468
Resilient worker	14.35 / 0.798	12.74 / 0.753	15.01 / 0.782	14.89 / 0.789	13.02 / 0.721	8.04 / 0.585
Fatigue-prone worker	10.10 / 0.655	8.72 / 0.620	7.46 / 0.529	7.91 / 0.521	8.72 / 0.579	1.43 / 0.382
8 tasks/day	11.85 / 0.931	8.94 / 0.816	13.06 / 0.973	13.12 / 0.974	12.51 / 0.934	6.84 / 0.723
16 tasks/day	9.00 / 0.505	9.09 / 0.518	7.32 / 0.417	7.02 / 0.405	9.30 / 0.509	1.70 / 0.331

S8 Reward-weight sweep

Table S5: AA-DQN as a function of the fatigue weight λ (attention weight $\mu = \lambda/2$; mean $\pm$ s.d. over three seeds).

λ ($\mu = \lambda/2$)	On-time rate	Mean fatigue	High-fatigue exposure	Break share
0.1	0.725 ± 0.008	0.324 ± 0.013	0.000	0.163